## → ♦ Apresentação

Curso: Machine Learning

Aluna: Camila Perazzo

Atividade de implementação de modelos e verificação de eficiência da estratégia adotada em problemas de regressão e classificação - 23/05/2023.

Este conjunto de dados contém classificações de consumo de combustível específicas do modelo e emissões estimadas de dióxido de carbono para novos veículos leves para venda no varejo no Canadá.

- MODELYEAR e.g. 2014
- MAKE e.g. Acura
- · MODEL e.g. ILX
- VEHICLE CLASS e.g. SUV
- ENGINE SIZE e.g. 4.7
- CYLINDERS e.g 6
- TRANSMISSION e.g. A6
- FUEL CONSUMPTION in CITY(L/100 km) e.g. 9.9
- FUEL CONSUMPTION in HWY (L/100 km) e.g. 8.9
- FUEL CONSUMPTION COMB (L/100 km) e.g. 9.2
- CO2 EMISSIONS (g/km) e.g. 182 --> low --> 0

## → ۞ Coleta e Preparação

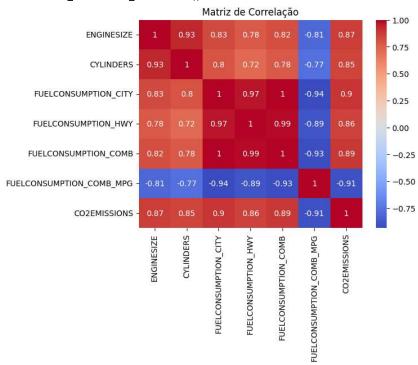
```
#Preparando o Ambiente
import pandas as pd
#Importando o arquivo de dados
df = pd.read_csv('/content/FuelConsumptionCo2.csv')
# Identifica os tipos de dados
df.dtypes
     MODELYEAR
                                   int64
     MAKE
                                  object
     MODEL
                                  object
     VEHICLECLASS
                                  object
     ENGINESIZE
                                  float64
     CYLINDERS
                                   int64
     TRANSMISSION
                                  object
     FUELTYPE
                                  object
     FUELCONSUMPTION CITY
                                  float64
     FUELCONSUMPTION HWY
                                  float64
     FUELCONSUMPTION COMB
                                  float64
     FUELCONSUMPTION_COMB_MPG
                                   int64
     CO2EMISSIONS
                                    int64
     dtype: object
#Verificando o formato do conjunto de dados
print("Shape do conjunto de dados: ",df.shape)
     Shape do conjunto de dados: (1067, 13)
#Visualizando amostra do conjunto de dados
df.head(10)
```

	MODELYEAR	MAKE	MODEL	VEHICLECLASS	ENGINESIZE	CYLINDERS	TRANSMISSION
0	2014	ACURA	ILX	COMPACT	2.0	4	AS5
1	2014	ACURA	ILX	COMPACT	2.4	4	M6
2	2014	ACURA	ILX HYBRID	COMPACT	1.5	4	AV7
3	2014	ACURA	MDX 4WD	SUV - SMALL	3.5	6	AS6
4	2014	ACURA	RDX AWD	SUV - SMALL	3.5	6	AS6
5	2014	ACURA	RLX	MID-SIZE	3.5	6	AS6

### 

```
2014 ACURA TLAWD
                                      MID-SIZE
                                                                             M6
df.columns
    'FUELCONSUMPTION_HWY', 'FUELCONSUMPTION_COMB',
'FUELCONSUMPTION_COMB_MPG', 'CO2EMISSIONS'],
          dtype='object')
import seaborn as sns
import matplotlib.pyplot as plt
# Selecionar apenas as colunas relevantes (excluindo a primeira coluna)
df_subset = df.iloc[:, 1:]
# Calcular a matriz de correlação
correlation_matrix = df_subset.corr()
# Exibir matriz de correlação com heatmap
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Matriz de Correlação')
plt.show()
```

<ipython-input-571-1a5a6818485c>:8: FutureWarning: The default value of numeric\_on
 correlation\_matrix = df\_subset.corr()



```
# Matrix de Correlação (Pearson)
corr = df.iloc[:, 1:].corr()
display(corr.style.background_gradient(cmap='BuGn').set_precision(2))
```

<ipython-input-572-06651845d749>:2: FutureWarning: The default value of numeric\_on
 corr = df.iloc[:, 1:].corr()

<ipython-input-572-06651845d749>:3: FutureWarning: this method is deprecated in fa display(corr.style.background\_gradient(cmap='BuGn').set\_precision(2))

	ENGINESIZE	CYLINDERS	FUELCONSUMPTION_CITY	FUEL(
ENGINESIZE	1.00	0.93	0.83	
CYLINDERS	0.93	1.00	0.80	
FUELCONSUMPTION_CITY	0.83	0.80	1.00	
FUELCONSUMPTION_HWY	0.78	0.72	0.97	
FUELCONSUMPTION_COMB	0.82	0.78	1.00	
FUELCONSUMPTION_COMB_MPG	-0.81	-0.77	-0.94	
CO2EMISSIONS	0.87	0.85	0.90	
4				<b>+</b>

### Atualize os nomes das colunas

	ENGINESIZE	CYLINDERS	FUELCONSCITY	FUELCONSHWY	FUELCONSCO
ENGINESIZE	1.00	0.93	0.83	0.78	0.
CYLINDERS	0.93	1.00	0.80	0.72	0.
FUELCONSCITY	0.83	0.80	1.00	0.97	1.
FUELCONSHWY	0.78	0.72	0.97	1.00	0.
FUELCONSCOMB	0.82	0.78	1.00	0.99	1.
FUELCONSCOMBMPG	-0.81	-0.77	-0.94	-0.89	-0.
CO2EMISSIONS	0.87	0.85	0.90	0.86	0.
4					<b>&gt;</b>

# 

```
# Criar novo DataFrame com as colunas desejadas
cdf = df.loc[:, ['ENGINESIZE', 'CYLINDERS', 'FUELCONSCOMB', 'CO2EMISSIONS']].copy()
cdf.head(10)
```

	ENG	SINESIZE	CYLINDERS	FUELCONSCOMB	CO2EMISSIONS	1
	0	2.0	4	8.5	196	
	1	2.4	4	9.6	221	
	2	1.5	4	5.9	136	
	3	3.5	6	11.1	255	
cdf.s	hape					
	(1067,	4)				
	6	3.5	6	10 1	232	

### 

4.1. Separe os dados de treino e teste

```
# bibliotecas utilizadas
import numpy as np
import warnings
# O Scikit-learn é uma biblioteca popular em Python para aprendizado de máquina e inclui o conjunto de dados Iris
from sklearn.linear_model import LogisticRegression
# Separando os dados de treino e teste
msk = np.random.rand(len(df)) < 0.80</pre>
dfTrain = cdf[msk]
dfTest = cdf[~msk]
print('Dataset de Treino = ', dfTrain.shape)
print('Dataset de Test = ', dfTest.shape)
     Dataset de Treino = (878, 4)
     Dataset de Test = (189, 4)
4.2. Import sklearn library
```

```
# Import sklearn library
from sklearn import linear_model
import seaborn as sns
import matplotlib.pyplot as plt
from scipy import stats
import numpy as np
```

4.3. Crie a variável com o modelo de Regressão Linear

```
# Linear Regression Model
regr = linear_model.LinearRegression()
```

4.4. Separe os dados de treino e teste

```
test_x_e = np.asanyarray(dfTest[['ENGINESIZE']])
test_x_c = np.asanyarray(dfTest[['CYLINDERS']])
test_x_f = np.asanyarray(dfTest[['FUELCONSCOMB']])
test_y = np.asanyarray(dfTest[['CO2EMISSIONS']])
```

4.5. Separe as Features de Treino e Teste

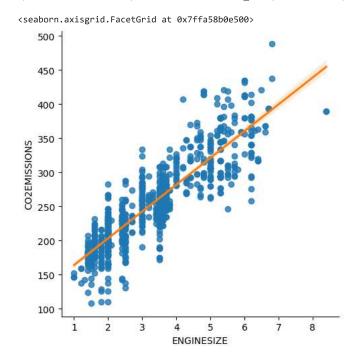
```
# Training X values (Variáveis independentes - Explanatory) e Y (Variável dependente - Target)
train_x_e = np.asanyarray(dfTrain[['ENGINESIZE']])
train_x_c = np.asanyarray(dfTrain[['CYLINDERS']])
train_x_f = np.asanyarray(dfTrain[['FUELCONSCOMB']])
train_y = np.asanyarray(dfTrain[['CO2EMISSIONS']])
```

4.6. Defina os valores para os coeficientes Theta 0: Intercepto e Theta 1: Coeficiente

```
# Ajuste da reta
{\tt regr.fit(train\_x\_e,\ train\_y)}
# Defining the values for coefficients - Theta 0: Intercepto e Theta 1: Coeficiente
print ('Coefficients: ', regr.coef_)
print ('Intercept: ', regr.intercept_)
     Coefficients: [[39.2397057]]
     Intercept: [124.9043427]
# Ajuste da reta
regr.fit(train\_x\_c,\ train\_y)
# Defining the values for coefficients - Theta 0: Intercepto e Theta 1: Coeficiente
print ('Coefficients: ', regr.coef_)
print ('Intercept: ', regr.intercept_)
     Coefficients: [[30.10877116]]
     Intercept: [81.60362912]
# Ajuste da reta
regr.fit(train_x_f, train_y)
# Defining the values for coefficients - Theta 0: Intercepto e Theta 1: Coeficiente
print ('Coefficients: ', regr.coef_)
print ('Intercept: ', regr.intercept_)
     Coefficients: [[15.89356536]]
     Intercept: [71.23987146]
```

### 4.7. Plot o Gráfico de Regressão Linear com a Reta Ajustada

sns.lmplot(x='ENGINESIZE', y='CO2EMISSIONS', line\_kws={"color": "C1"}, data=dfTrain)



 $sns.lmplot(x='CYLINDERS', y='CO2EMISSIONS', line\_kws=\{"color": "C2"\}, data=dfTrain)$ 

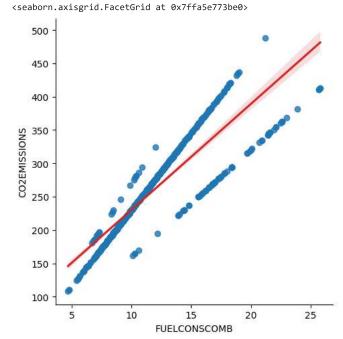
```
<seaborn.axisgrid.FacetGrid at 0x7ffa5e7e89d0>

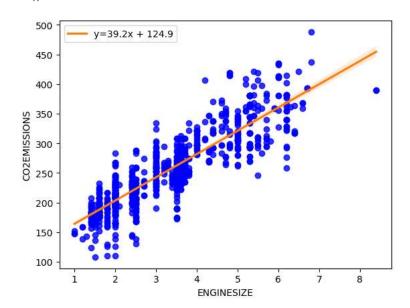
500 -

450 -

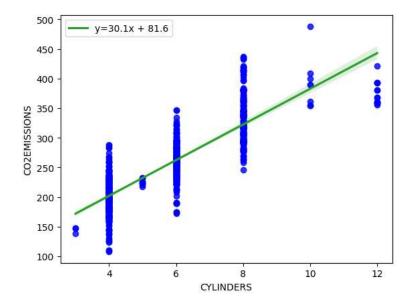
400 -

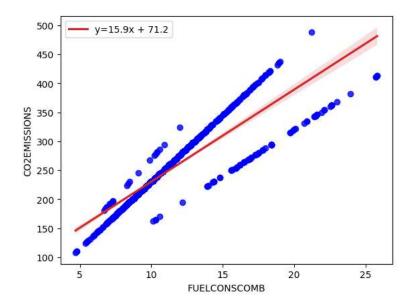
sns.lmplot(x='FUELCONSCOMB', y='CO2EMISSIONS', line_kws={"color": "C3"}, data=dfTrain)
```





```
ax.legend()
plt.show()
```





### 4.8. Encontre o R2-Score

```
import sklearn
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression

# Combinações de variáveis
combinations = [('FUELCONSCOMB',), ('CYLINDERS',), ('ENGINESIZE',)]

best_r2_score = 0
best_variable_combination = None

for combination in combinations:
    # Separar as variáveis independentes (X) e a variável dependente (y)
    X = df[list(combination)]
    y = df['CO2EMISSIONS']
```

```
# Separar os dados em conjunto de treinamento e teste
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
# Criar o modelo de Regressão Linear
regression_model = LinearRegression()
# Treinar o modelo usando o conjunto de treinamento
regression_model.fit(X_train, y_train)
# Calcular o R2-Score
r2 score = regression model.score(X test, y test)
print("Variáveis:", combination)
print("R2-Score:", r2_score)
Variáveis: ('FUELCONSCOMB',)
 R2-Score: 0.8071474868274242
 Variáveis: ('CYLINDERS',)
 R2-Score: 0.7317140029783895
Variáveis: ('ENGINESIZE',)
R2-Score: 0.7615595731934373
```

### © Encontre o Melhor ajuste (R2-Score)

Entre as seguintes combinações de variáveis:

- FUELCONSCOMB & CO2EMISSIONS
- CYLINDERS & CO2EMISSIONS
- ENGINESIZE & CO2EMISSIONS

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
# Lista das combinações de variáveis
combinations_pair = [('FUELCONSCOMB', 'CO2EMISSIONS'),
               ('CYLINDERS', 'CO2EMISSIONS'),
               ('ENGINESIZE', 'CO2EMISSIONS')]
best_r2_score = -1 # Inicialização do melhor R2-Score
for feature, target in combinations_pair:
    # Separe os dados de treino e teste
    train_cdf, test_cdf = train_test_split(cdf, test_size=0.2, random_state=42)
    # Crie a variável com o modelo de Regressão Linear
    model = LinearRegression()
    # Separe as Features de Treino e Teste
    train_x = train_cdf[[feature]]
    train_y = train_cdf[target]
    test_x = test_cdf[[feature]]
    test_y = test_cdf[target]
    # Rode o Modelo
    model.fit(train_x, train_y)
    # Encontre o R2-Score
    y_pred = model.predict(test_x)
    r2 = r2_score(test_y, y_pred)
    if r2 > best_r2_score:
        best_r2_score = r2
        best_feature = feature
        best_target = target
print("Melhor ajuste:")
print("Variável de Feature:", best_feature)
print("Variável de Target:", best_target)
print("Melhor R2-Score:", best_r2_score)
     Melhor ajuste:
     Variável de Feature: FUELCONSCOMB
     Variável de Target: CO2EMISSIONS
     Melhor R2-Score: 0.8071474868274242
```

# Quality Regressão Linear Múltipla

Usando as variáveis:

- ENGINESIZE
- CYLINDERS
- FUELCONSCOMB

```
# Import sklearn library
from sklearn import linear model
import seaborn as sns
import matplotlib.pyplot as plt
from scipy import stats
# Linear Regression Model
regr = linear_model.LinearRegression()
# Training X values (Var Independent - Explanatory) e Y (Var Dependent - Target)
train_x = np.asanyarray(dfTrain[['ENGINESIZE', 'CYLINDERS', 'FUELCONSCOMB']])
train_y = np.asanyarray(dfTrain[['CO2EMISSIONS']])
# Performing Line Adjustment
regr.fit(train_x, train_y)
# Defining the values for coefficients - Theta 0: Intercepto e Theta 1: Coeficiente
print ('Coefficients: ', regr.coef_)
print ('Intercept: ', regr.intercept_)
     Coefficients: [[11.43455109 7.3905759 9.28830002]]
     Intercept: [67.05954605]
```